

SQL Queries

Joins over Two or More Tables

SQL Queries

```
SELECT [DISTINCT] [SUM | COUNT | AVG] result_table  
FROM table1, table2  
[WHERE table_predicates AND join_conditions]  
[GROUP BY grouping_attributes  
    [HAVING agg_condition]]  
[ORDER BY sorting_attributes]  
[UNION [ALL]] [INTERSECT] [EXCEPT]
```

Cartesian Product

- $R(A) = \{1,1,2,3\}$
- $S(B) = \{1,3,4\}$
- $R \times S(A,B) = \{(1,1), (1,3), (1,4), (1,1), (1,3), (1,4), (2,1), (2,3), (2,4), (3,1), (3,3), (3,4)\}$
- The result consists of pairs of one element from R and one from S
- Every element from R is paired with every element from S
- The number of elements in $R \times S$ is $|R|*|S|$, i.e., the size of R multiplied by the size of S
- select *
from R, S
- The schema of the result is the **union** of the R schema and the S schema

Cartesian Product Generalization

- $R(A) = \{1,1,2,3\}$
- $S(B) = \{1,3,4\}$
- $T(C) = \{2,4\}$
- $R \times S(A,B) = \{(1,1),(1,3),(1,4), (1,1),(1,3),(1,4), (2,1),(2,3),(2,4), (3,1),(3,3),(3,4)\}$
- select * from R, S
- $R \times S \times T(A,B,C) = \{(1,1,2),(1,3,2),(1,4,2), (1,1,2),(1,3,2),(1,4,2), (2,1,2),(2,3,2),(2,4,2), (3,1,2),(3,3,2),(3,4,2), (1,1,4),(1,3,4),(1,4,4), (1,1,4),(1,3,4),(1,4,4), (2,1,4),(2,3,4),(2,4,4), (3,1,4),(3,3,4),(3,4,4)\}$
- select * from R, S, T

Two-Table Join

- $R(A) = \{1,1,2,3\}$
- $S(B) = \{1,3,4\}$
- $R \bowtie_{A=B} S = \{$
 $\underline{(1,1)}, \underline{(1,3)}, \underline{(1,4)},$
 $\underline{(1,1)}, \underline{(1,3)}, \underline{(1,4)},$
 $\underline{(2,1)}, \underline{(2,3)}, \underline{(2,4)},$
 $\underline{(3,1)}, \underline{(3,3)}, \underline{(3,4)}\} = \{(1,1), (1,1), (3,3)\}$
- Join condition between attributes from the two tables
- Only those tuples from the Cartesian product that satisfy the join condition are included in the result

- select * from R, S
where **A = B**
- Condition does not have to be equality
- select * from R, S
where **A > B**
 - $\{(2,1), (3,1)\}$

Multiple-Table Join

- $R(A) = \{1,1,2,3\}$
- $S(B) = \{1,3,4\}$
- $T(C) = \{2,4\}$
- select * from R, S, T
where **A=B and B>C**
- If there is no condition for a table, Cartesian product is performed for that table
- $R \bowtie_{A=B} S \bowtie_{B>C} T(A,B,C) = \{$
~~(1,1,2),(1,3,2),(1,4,2),~~
~~(1,1,2),(1,3,2),(1,4,2),~~
~~(2,1,2),(2,3,2),(2,4,2),~~
~~(3,1,2),(3,3,2),(3,4,2),~~
~~(1,1,4),(1,3,4),(1,4,4),~~
~~(1,1,4),(1,3,4),(1,4,4),~~
~~(2,1,4),(2,3,4),(2,4,4),~~
~~(3,1,4),(3,3,4),(3,4,4)\} = \{(3,3,2)\}~~

Duplicate Attribute Names

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- select * from Product, PC
 - schema: (maker, **Product.model**, type, **PC.model**, speed, ram, hd, price)
 - select **Product.model**, maker, price from Product, PC
 - select **P.model**, maker, PC.price from **Product P**, PC

Join Query Examples

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- select * from Product P, PC
where **P.model = PC.model**
- select P1.maker, PC.model AS pc_model, L.model AS laptop_model
from **Product P1, Product P2**, PC, Laptop L
where **P1.maker = P2.maker and P1.model = PC.model and P2.model = L.model and PC.price > L.price**
 - Find the (PCs, laptop) pairs produced by the same maker for which the PC price is larger than the laptop price
 - Multiple instances of a table can appear in a query. They have to be renamed as the attributes are renamed.

Abstract Evaluation Model

- select P1.maker, PC.model AS pc_model, L.model AS laptop_model
from Product P1, Product P2, PC, Laptop L
where P1.maker = P2.maker and P1.model = PC.model and
P2.model = L.model and PC.price > L.price
- **For** each tuple P1 in table Product
 - For** each tuple P2 in table Product
 - For** each tuple PC in table PC
 - For** each tuple L in table Laptop
 - if** P1.maker = P2.maker and P1.model = PC.model and P2.model =
L.model and PC.price > L.price
 - then** add(P1.maker, PC.model, L.model) to the result

Abstract Evaluation Model for General Queries

SELECT [DISTINCT] [SUM | COUNT | AVG] result_table

FROM table₁, table₂, ...

[WHERE table_predicates AND join_conditions]

[GROUP BY grouping_attributes

[HAVING agg_condition]]

[ORDER BY sorting_attributes]

[UNION [ALL]] [INTERSECT] [EXCEPT]

- **The evaluation model for joins is first applied to the entire WHERE clause**
- **Everything else is evaluated on the result of the join evaluation**

Examples

- Computers
- TPC-H